

Asymmetric Penalty in Multi-Branch Neural Networks for Non-Stationary Hydrological Forecasting

Tan Khoi Nguyen Nguyen¹[0009-0009-7970-9883], Thu Le²[0009-0008-3480-8311], Thai Tran Quang³[0009-0000-4674-4388], Khai Trinh Tien⁴[0009-0002-0447-5743], and Phuoc Phan Tan⁵[0009-0003-7078-9406]

FPT University, Ho Chi Minh Campus, Vietnam
nguyennnguyentankhoi@gmail.com, thulvm@fe.edu.vn, {tranquangthai1029, khaitrinh2005, phanphuoc9113}@gmail.com

Abstract. Monitoring electrical conductivity (EC) is vital for managing salinity intrusion, yet current Long-term Time-series Forecasting (LTSF) models often fail to capture sudden, extreme spikes. This limitation stems from standard symmetric loss functions like Mean Squared Error (MSE), which penalize all errors equally and bias predictions toward the conditional mean, resulting in dangerous "oversmoothing" of hazardous peaks. To address this, we propose a hybrid decomposition-ensemble network designed for peak sensitivity. Our approach employs CEEMDAN to decouple the EC signal into noisy high-frequency spikes and deterministic trends, routed respectively into deep MLP and linear branches. Crucially, we introduce an Asymmetric Spike-Aware Loss that heavily penalizes underestimations, forcing the network to respect the true magnitude of fluctuations. Evaluated against baselines like PatchTST and CEEMDAN-LTSF-Linear on USGS sensor data (2021–2025), our model effectively captures extreme contamination events that standard methods overlook. This framework provides a robust early-warning mechanism for critical hydrological monitoring.

Keywords: Water quality forecasting, CEEMDAN, Multi-Layer Perceptron, Asymmetric Loss, Time-series analysis

1 Introduction

Monitoring water quality in aquacultural and riverine systems relies heavily on predicting Electrical Conductivity (EC), a primary proxy for salinity intrusion exacerbated by climate change [8, 13, 4]. As hydrological patterns become more erratic, high-frequency EC data is increasingly non-stationary and dominated by sudden, high-magnitude fluctuations, or "spikes". Accurately forecasting these extreme events is critical for early-warning systems but remains remarkably difficult for conventional time-series models.

Recently, Long-term Time-series Forecasting (LTSF) has seen rapid advancements, ranging from complex Transformer architectures [10] to simpler, highly stable linear mappings [11, 3]. While these deep learning models excel at capturing general seasonal trends, they share a critical blind spot when evaluated on highly volatile environmental data: they tend to "smooth out" the spikes. This occurs because most models are optimized using symmetric loss functions, typically Mean Squared Error (MSE), which penalize overestimations and underestimations equally. To minimize the global error, the network converges toward the conditional mean, systematically underestimating hazardous peaks and rendering the forecast practically useless during an actual crisis.




Fig. 1. The smoothing effect of traditional symmetric-loss optimization during a high-magnitude EC spike

To address this critical underestimation, forecasting models must separate noise from the underlying trend and heavily penalize under-predictions. In this paper, we propose a Hybrid Decomposition-Ensemble model explicitly designed to balance trend stability with peak sensitivity. By applying CEEMDAN directly to the raw signal, the decoupled components are routed into a multi-branch network: a deep Multi-Layer Perceptron (MLP) tackles high-frequency stochastic spikes, while a linear branch processes low-frequency deterministic trends. Crucially, we replace MSE with an Asymmetric Spike-Aware Loss that imposes a heavy penalty whenever the network misses a high-magnitude peak.

The primary contributions of this study are summarized as follows:

1. **Frequency-Decoupled Architecture:** We propose a specialized multi-branch network that models stochastic spikes and deterministic trends independently, significantly reducing the interference between high-frequency volatility and seasonal patterns.
2. **Asymmetric Optimization Strategy:** We introduce a custom loss function combining focal weighting and directed asymmetric penalties. This explic-

itly counteracts the dangerous smoothing bias caused by standard MSE optimization in volatile time-series.

3. Empirical Validation on Raw Scales: Extensive experiments on high-resolution USGS datasets (2021-2025) demonstrate that the proposed method achieves competitive performance against state-of-the-art baselines, offering vastly improved accuracy in capturing the morphology of extreme water quality events.

2 Related Work

2.1 Deep Learning and Hybrid Forecasting in Hydrology

Water quality forecasting has shifted from traditional statistical methods (e.g., ARIMA, SVR) to deep learning to better capture the non-linear, volatile nature of hydrological data, with RNNs (LSTM, GRU) becoming the standard.

To enhance predictive accuracy, researchers often pair these networks with signal decomposition techniques. For example, coupling CEEMDAN [9] with LSTM [8, 5] significantly reduces background noise and lowers the root mean square error (RMSE) for indicators like ammonia nitrogen. Similarly, LTSF-Linear combined with CEEMDAN has been utilized to manage seasonal variations in water treatment [7]. However, these noise-filtering hybrid pipelines predominantly rely on symmetric loss functions (e.g., MSE) during training. This inadvertently biases models toward the mean, leading to a systematic underestimation of extreme contamination events.

2.2 The LTSF Architecture Debate

In the broader context of Long-term Time-series Forecasting (LTSF), the dominance of Transformers has been recently challenged. Zeng et al. [11] demonstrated that for temporal data with limited semantic density, simple linear mappings (like DLinear) often outperform complex attention mechanisms by avoiding temporal noise overfitting. Building on this, PatchTST [6] showed that dividing time-series into patches captures local semantic patterns more effectively than point-wise processing.

Our work bridges these findings through a hybrid strategy. We leverage the proven stability of linear mappings for low-frequency deterministic trends, while reserving the non-linear capacity of Deep Multi-Layer Perceptrons (MLPs)—which have recently shown a resurgence in time-series forecasting [1]—strictly for high-frequency stochastic spikes.

3 Proposed Methodology

The proposed forecasting framework is explicitly designed to address the non-stationarity and volatile "spikes" inherent in high-frequency Electrical Conductivity (EC) data. Rather than forcing a single neural architecture to learn the entire complex signal spectrum, we adopt a "divide-and-conquer" strategy through

a three-phase sequential pipeline: Adaptive Signal Decomposition, Multi-Branch Hybrid Architecture, and Asymmetric Optimization.




Fig. 2. The pipeline of the proposed hybrid decomposition-ensemble method

3.1 Adaptive Signal Decomposition

Raw environmental time-series data is typically a superimposition of long-term seasonal trends and abrupt, high-frequency stochastic fluctuations (e.g., sudden salinity intrusions). Feeding such raw sequences directly into standard deep learning models often results in gradient noise and poor convergence. To mitigate this, we employ the Complete Ensemble Empirical Mode Decomposition with Adaptive Noise (CEEMDAN) algorithm to decouple the original time series $X(t)$ into K Intrinsic Mode Functions (IMFs) and a residual trend $R(t)$. The decomposition process is mathematically expressed as:

$$X(t) = \sum_{k=1}^K IMF_k(t) + R(t) \quad (1)$$

A critical distinction of our approach is that we apply CEEMDAN directly to the unscaled, raw signal. Conventional pipelines typically normalize the entire dataset before decomposition, which inadvertently compresses the magnitude of extreme EC peaks. By retaining the original scale, we preserve the morphological fidelity of the critical spikes.

Setting $K = 11$, we partition the resulting components into two distinct frequency-based sets:

1. High-Frequency Set (X_{high}): Comprising IMF_1 through IMF_3 , this set isolates rapid fluctuations and stochastic noise.
2. Low-Frequency Set (X_{low}): Comprising IMF_4 through IMF_{11} alongside the residue $R(t)$, this set represents the underlying deterministic trends.




Fig. 3. CEEMDAN decomposition results. High-frequency components (IMFs 1-3, in red) isolate sharp stochastic spikes, whereas low-frequency components (IMFs 4-11 and residue, in blue) capture smooth deterministic trends

3.2 Hybrid Network Architecture

Instead of employing a generalized recurrent or attention-based network, we route the decoupled frequency sets into specialized neural branches. This architecture significantly reduces signal interference between short-term volatility and long-term periodicity.

1. **Deep MLP Branch:** Because EC spikes exhibit high non-linearity, we utilize a Deep Multi-Layer Perceptron (MLP) featuring wide hidden layers (512 units) and ReLU activations to act as a dedicated spike detector for the high-frequency components.
2. **Trend Linear Branch:** Inspired by the robust performance of DLinear models, the low-frequency components are processed by a single linear mapping layer. This design choice prevents overfitting to minor temporal fluctuations, ensuring stable extrapolation of macroscopic trends.
3. **Auxiliary MLP:** A compact MLP (154 units) is employed to encode the supplementary statistical features.

The latent representations from the three branches are concatenated and passed through a Fusion Layer. Crucially, the network predicts a temporal delta ($\Delta\hat{y}$) rather than absolute values. The final forecast is reconstructed by adding this delta to the last observed value x_t :

$$\hat{y}_{t+H} = x_t + \text{Fusion}(y_{high}, y_{low}, y_{feat}) \quad (2)$$

3.3 Asymmetric Optimization Strategy

The core technical contribution of this study is the optimization strategy. Baseline models (e.g., PatchTST, CEEMDAN-LTSF-Linear) systematically fail to

capture high-magnitude spikes due to their reliance on symmetric loss functions, typically Mean Squared Error (MSE). Because MSE penalizes overestimations and underestimations equally, the network’s weights are forced to converge toward the conditional mean to minimize global error. This mean-reverting bias results in the dangerous "oversmoothing" of extreme risk events.

Inspired by traditional asymmetric penalty theories in risk management [2] and recent generalized extreme value loss strategies [12], we propose the Asymmetric Spike-Aware Loss ($\mathcal{L}_{SA}$) to enforce peak sensitivity:

$$\mathcal{L}_{SA} = \frac{1}{N} \sum_{i=1}^N \omega_i \cdot (1 - e^{-|y_i - \hat{y}_i|})^\gamma \cdot (y_i - \hat{y}_i)^2 \quad (3)$$

Where the asymmetric weight ω_i introduces a directed penalty, forcing the model to heavily penalize "False Negatives" (underestimating an actual spike):

$$\omega_i = \begin{cases} \alpha & \text{if } (y_i - \hat{y}_i) > 0 \\ 1.0 & \text{otherwise} \end{cases} \quad (4)$$

By empirically setting $\alpha = 5.0$, the loss function penalizes under-predictions five times more severely than over-predictions. The focal parameter $\gamma = 2.0$ dynamically down-weights gradients from minor errors during stable water periods, allowing the network to concentrate its computational capacity on learning abrupt shocks.

4 Experimental Setup

4.1 Dataset Description and Partitioning

To evaluate the framework, we utilize an hourly water quality dataset (43,813 observations) from the USGS sensor network, spanning January 2021 to December 2025. The primary target variable is Electrical Conductivity (EC), a critical proxy for salinity intrusion. To prevent data leakage inherent in temporal dependencies, we employ a strict chronological split: 60% for training (2021–2023), 20% for validation (2024), and 20% for testing (2025).

Table 1. Descriptive statistics of the USGS Electrical Conductivity dataset.

Statistic	Value
Time Period	Jan 2021 – Dec 2025
Total Observations	43,813
Unit	$\mu\text{S}/\text{cm}$
Mean	202.123
Standard Deviation (σ)	43.846
Minimum	68.0
Maximum (Peak)	440.0

4.2 Baseline Models

To rigorously assess the superiority of the proposed framework, we benchmark it against four state-of-the-art forecasting models representing different architectural paradigms:

1. Transformer: The standard self-attention mechanism, representing deep global semantic modeling.
2. PatchTST: A state-of-the-art Transformer variant that utilizes patching and channel independence to capture local semantics.
3. CEEMDAN-LTSF-Linear: A highly efficient linear mapping model (specifically DLinear) that excels at extracting global trends without overfitting temporal noise.
4. CEEMDAN-LSTM: A classic hybrid approach combining signal decomposition with recurrent networks, optimized via standard symmetric loss.

4.3 Implementation Details

All models were implemented in PyTorch on an NVIDIA RTX 4060 GPU. The historical look-back window (L) is fixed at 168 hours to capture weekly seasonality, forecasting future horizons $T \in \{6, 12, 24, 48, 96, 168\}$ hours. Our proposed model uses CEEMDAN configured with $K = 11$ modes and 20 ensemble trials. Training utilizes the Adam optimizer (initial learning rate 10^{-3} , batch size 64) with a 10-epoch early stopping patience. The Asymmetric Spike-Aware Loss parameters are empirically set to $\alpha = 5.0$ and $\gamma = 2.0$. Baselines retain their default optimal configurations.

4.4 Evaluation Metrics

To quantify forecasting performance, we rely on three established accuracy measures: Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and the Coefficient of Determination (R^2). To specifically evaluate the model’s capability in early-warning systems, we introduce two domain-specific metrics. Peak MAE calculates the mean absolute error strictly for observations exceeding a critical hazard threshold (set at $250 \mu S/cm$). Spike Recall measures the percentage of these true extreme events that the model successfully predicts above the threshold.

5 Results and Discussion

5.1 Quantitative Performance Analysis

Tables 2 and 3 show that standard baselines, particularly CEEMDAN-LTSF-Linear, yield better global metrics (MAE, RMSE) than our proposed model across most horizons. For instance, at 168 hours, CEEMDAN-LTSF-Linear achieves an RMSE of 12.442 ($R^2 = 0.939$) compared to our 21.841.

However, this numerical gap is the intended outcome of our asymmetric optimization strategy. Standard models trained with symmetric losses (e.g., MSE) achieve low global errors by over-fitting the stable majority of the data, effectively treating rare but critical salinity spikes as outliers. This results in a low-pass filter behavior that is numerically superior on paper but practically dangerous for real-world extreme event prediction.

Table 2. Short-term Forecasting Performance Comparison (Horizon $\tau \in \{6, 12, 24\}$ Hours)

Model	Horizon = 6 Hours			Horizon = 12 Hours			Horizon = 24 Hours		
	MAE	RMSE	R^2	MAE	RMSE	R^2	MAE	RMSE	R^2
Transformer	6.065	7.644	0.977	3.754	6.135	0.985	5.413	8.763	0.970
PatchTST	3.502	5.909	0.986	3.485	6.174	0.985	4.945	8.171	0.974
CEEMDAN-LSTM	2.904	5.317	0.989	3.080	5.789	0.987	4.781	7.733	0.976
CEEMDAN-LTSF-Linear	1.943	4.125	0.993	2.135	4.581	0.992	2.648	5.341	0.989
Proposed Model	4.561	7.047	0.980	5.729	8.581	0.971	7.202	10.668	0.954

Table 3. Long-term Forecasting Performance Comparison (Horizon $\tau \in \{48, 96, 168\}$ Hours)

Model	Horizon = 48 Hours			Horizon = 96 Hours			Horizon = 168 Hours		
	MAE	RMSE	R^2	MAE	RMSE	R^2	MAE	RMSE	R^2
Transformer	8.640	12.653	0.937	12.793	17.463	0.879	17.414	22.816	0.792
PatchTST	7.896	11.711	0.946	11.296	15.740	0.902	15.884	21.533	0.814
CEEMDAN-LSTM	7.406	11.489	0.948	10.552	15.209	0.908	14.653	20.831	0.826
CEEMDAN-LTSF-Linear	3.821	6.891	0.981	5.825	9.561	0.964	8.075	12.442	0.939
Proposed Model	9.349	14.009	0.922	12.134	17.289	0.881	15.393	21.841	0.811

5.2 The Cost of False Negatives in Salinity Monitoring

Table 4 quantifies the operational advantage of our asymmetric optimization strategy during actual extreme events ($EC > 250 \mu\text{S/cm}$). While aggregate metrics (Tables 2 and 3) favor symmetric-loss baselines, evaluating strictly on hazard zones reveals a different reality. As the forecasting horizon extends, the CEEMDAN-LTSF-Linear baseline suffers from severe oversmoothing. By $T=168$, its Peak MAE degrades to 12.894, missing nearly 17% of critical surges. In contrast, driven by the $\mathcal{L}_{SA}$ penalty, our proposed model preserves the true magnitude of fluctuations. At $T=168$, it maintains a Peak MAE of 6.454—halving the baseline’s error—while successfully capturing 99.4% of extreme events. This confirms that sacrificing precision during calm periods for sustained peak sensitivity is an operational necessity for robust early-warning systems.

Table 4. Extreme Event Forecasting Performance ($EC > 250 \mu S/cm$)

Horizon Total Spikes		CEEMDAN-LTSF-Linear		Proposed Model	
		Peak MAE	Recall (%)	Peak MAE	Recall (%)
T=6	2,473	2.414	98.30	3.868	99.30
T=12	2,472	3.179	97.49	3.460	98.10
T=24	2,472	4.890	96.85	3.522	98.90
T=48	2,471	8.076	93.58	3.926	98.50
T=96	2,471	11.083	89.74	4.728	98.10
T=168	2,471	12.894	83.39	6.454	99.40

5.3 Visual Case Study




Fig. 4. Model predictions during a critical salinity spike. Baselines (a-d) oversmooth the peak amplitude, whereas the Proposed Model (e) accurately captures the extreme event.

To prove that our model excels where standard metrics fail, we isolated a specific high-magnitude salinity intrusion event for visual comparison.

In contrast, our proposed model captures the sharp morphology of the event. Even though its predictions fluctuate more during the pre-spike period (which inflates its MAE), it successfully traces the peak when it matters most. This confirms that decoupling the signal via CEEMDAN and applying an asymmetric penalty successfully forces the network to act as an extreme event detector rather than a simple trend extrapolator.

5.4 Model Interpretability and Error Distribution

We applied SHAP to see exactly how the multi-branch architecture makes decisions during severe salinity intrusions. Since the raw SHAP values are extracted directly from the network’s normalized Δ predictions, their numerical scale is quite small. To make the comparison easier to follow, we mapped these absolute scores into relative percentages (Figure 5). When extreme spikes hit, the High-Frequency branch (IMFs 1-3) clearly takes over. It drives the forecast by accounting for 46.3% of the variance. This lines up well with our core idea: the best way to catch sudden environmental shocks is by isolating high-frequency noise. Meanwhile, the Low-Frequency branch is kept in check, contributing just 34.3%. We specifically constrained it this way so that stable macroscopic trends wouldn’t drag down the amplitude of sudden spikes. The remaining 19.5% comes from the Auxiliary branch, which mostly works in the background to provide contextual stability.




Fig. 5. Branch importance analysis using SHAP values. The High-Frequency branch dominates the decision-making process during extreme salinity intrusions with the highest impact score.

The practical consequence of this signal-decoupled optimization is visually confirmed through the error distribution analysis in Figure 6. The global best-performing baseline, CEEMDAN-LTSF-Linear, tracks actual values adequately in normal ranges but exhibits a severe ceiling effect when Electrical Conductivity

exceeds $250 \mu\text{S}/\text{cm}$. Driven by the Asymmetric Spike-Aware Loss ($\mathcal{L}_{SA}$), the Proposed Model effectively mitigates this under-prediction bias. By sacrificing tight variance at the lower end, it maintains alignment with the ideal $y = x$ reference line well beyond the $250 \mu\text{S}/\text{cm}$ threshold.




Fig. 6. Tail-behavior analysis via scatter plot of predictions versus ground truth. CEEMDAN-LTSF-Linear suffers from a severe ceiling effect above $250 \mu\text{S}/\text{cm}$, whereas the proposed model maintains alignment with the ideal $y = x$ reference line.

6 Conclusion

This paper presented a novel hybrid framework to overcome the oversmoothing limitations of standard LTSF models in forecasting high-frequency salinity intrusions. Relying on CEEMDAN decomposition, a multi-branch architecture, and an Asymmetric Spike-Aware Loss ($\mathcal{L}_{SA}$), our method shifts the paradigm from aggregate error minimization to extreme event detection. Evaluations demonstrate the model’s unique capability to maintain morphological fidelity during extreme surges ($>250 \mu\text{S}/\text{cm}$), effectively eliminating the ceiling effect of baselines. SHAP analysis further confirms that the network successfully prioritizes high-frequency stochastic noise during critical events. Despite these advantages, the framework introduces certain trade-offs. The strict penalty on under-predictions naturally increases variance during stable periods, which could occasionally trigger false positive alerts. Additionally, applying CEEMDAN directly to raw signals entails higher computational overhead. Future research will explore lightweight signal decomposition techniques and automated hyperparameter optimization (e.g., for α and γ) to enhance the model’s efficiency and adaptability across diverse datasets. Ultimately, this research provides an interpretable and operationally viable early-warning mechanism for estuarine agriculture.

References

1. Chen, S.A., Li, C.L., Yoder, N., Arik, S.O., Pfister, T.: TSMixer: An all-MLP architecture for time series forecasting. *Transactions on Machine Learning Research* (2023)
2. Christoffersen, P.F., Diebold, F.X.: Optimal prediction under asymmetric loss. *Econometric Theory* **13**(6), 808–817 (1997)
3. Li, Z., Qi, S., Li, Y., Xu, Z.: Revisiting long-term time series forecasting: An investigation on linear mapping. *arXiv preprint arXiv:2305.10721* (2023)
4. Lin, F., Li, X., Su, Y., Yan, J.: Water quality prediction model based on improved long short-term memory neural network and empirical mode decomposition. *Discover Artificial Intelligence* **5**, 1–18 (2025). <https://doi.org/10.1007/s44163-025-00454-y>
5. Luo, X., et al.: A hybrid model for water quality prediction based on CEEMDAN and deep learning. *Environmental Science and Pollution Research* **29**, 45472–45486 (2022)
6. Nie, Y., Nguyen, N.H., Sinthong, P., Kalagnanam, J.: A time series is worth 64 words: Long-term forecasting with Transformers. In: *The Eleventh International Conference on Learning Representations (ICLR)* (2023)
7. Shin, J., Yoon, S., Park, N.S., Kim, Y.: Development of water quality prediction model using LTSF-Linear and complete ensemble empirical mode decomposition. *Desalination and Water Treatment* **323**, 101254 (2025). <https://doi.org/10.1016/j.dwt.2025.101254>
8. Su, Y.: Research on water quality prediction in Shanghai based on CEEMDAN-LSTM model. In: *Proceedings of the 2023 7th International Conference on Electronic Information Technology and Computer Engineering*. pp. 945–951. Association for Computing Machinery, New York, NY, USA (2024). <https://doi.org/10.1145/3650400.3650560>
9. Torres, M.E., Colominas, M.A., Schlotthauer, G., Flandrin, P.: A complete ensemble empirical mode decomposition with adaptive noise. In: *2011 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. pp. 4144–4147. IEEE (2011)
10. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, Ł., Polosukhin, I.: Attention is all you need. In: *Advances in Neural Information Processing Systems*. vol. 30 (2017)
11. Zeng, A., Chen, M., Zhang, L., Xu, Q.: Are Transformers effective for time series forecasting? In: *Proceedings of the AAAI Conference on Artificial Intelligence*. vol. 37, pp. 11121–11128 (2023). <https://doi.org/10.1609/aaai.v37i9.26317>
12. Zhang, M., Ding, D., Pan, X., Yang, M.: Enhancing time series predictors with generalized extreme value loss. *IEEE Transactions on Knowledge and Data Engineering* **35**(2), 1473–1487 (2023). <https://doi.org/10.1109/TKDE.2021.3108831>
13. Zhu, Q., Yu, X., Zhao, L., Xu, L.: A hybrid deep learning model for water quality prediction. *IEEE Access* **13**, 106406–106415 (2025). <https://doi.org/10.1109/ACCESS.2025.3580741>