Spatiotemporal Housing Price Prediction: Comparing Functional Semiparametric Models with Modern Machine Learning Approaches

Nguyen Trinh Quy[1], Nguyen Cong Minh[1], Le Pham Hue Chi[1], and Nguyen Do Nhat Quang[1]

FPT University HCM,
Ho Chi Minh City, Vietnam

Abstract. Accurate forecasting of multi-region residential house prices is a central challenge in real-estate economics with implications for policy, investment, and urban planning. This paper extends the functional spatiotemporal semiparametric mixed-effects (FST-SM) model proposed by Chen and Zheng [1] by systematically comparing it against four modern machine learning (ML) approaches — Linear Regression, Random Forest, XGBoost, LightGBM, and a Long Short-Term Memory (LSTM) network — applied to the Zillow Home Value Index (ZHVI) across 894 U.S. metropolitan statistical areas from January 2001 to January 2026, comprising 268 200 region-month observations. We target monthly log-returns (a stationary, scale-free transformation of ZHVI) and engineer autoregressive lag features (lag-1 through lag-12) together with a city-level target encoding. A temporal 80/20 train/test split ensures strict non-leakage evaluation on future data. Experimental results show that Linear Regression achieves the lowest log-return RMSE (0.00276) and highest R^2 (0.876), outperforming Random Forest, XGBoost, and LightGBM, while LSTM performs worst (RMSE 0.00515, R^2 0.569) despite its temporal modelling capacity. Feature importance analysis reveals that the one-month lag dominates all tree-based models, echoing the FST-SM finding that lag-1 carries the strongest fixed effect ($\hat{\beta}_1 = 2.39$). Our results suggest that the dominant temporal structure of ZHVI is captured effectively by linear autoregression, and that the additional inductive bias of the FST-SM's functional principal component decomposition and CAR spatial structure provide competitive advantages particularly for local community-level patterns, whereas scalable ML pipelines are more practical for national-scale forecasting tasks. All code and processed data are released to support reproducibility.

Keywords: Housing price prediction · Zillow ZHVI · Functional mixed-effects model · Machine learning · Gradient boosting · LSTM · Spatiotemporal analysis

1 Introduction

Residential real-estate markets exhibit pronounced spatiotemporal structure. Prices in geographically proximate markets tend to co-move, while long-run tra-

jectories diverge substantially across cities and communities. Accurate, scalable prediction of these dynamics is economically significant: it informs mortgage risk assessment, zoning decisions, urban investment, and household financial planning [12].

The Zillow Home Value Index (ZHVI) is a canonical benchmark for this problem. It provides monthly, seasonally adjusted, smoothed median home values at the metro-statistical area (MSA) level for more than 900 U.S. cities, covering over two decades of market fluctuations including the 2007–2012 financial crisis, the COVID-19 pandemic price surge of 2020–2022, and the subsequent correction [2].

Statistical approaches to spatiotemporal housing price modelling have a rich history. Hedonic regression [3] accounts for property attribute heterogeneity; spatial-lag models [4] capture geographic autocorrelation; and functional data analysis (FDA) methods [5] represent price trajectories as continuous curves. The recent FST-SM model of Chen and Zheng [1] elegantly unifies these threads: it decomposes ZHVI trajectories via functional principal component analysis (FPCA), models spatial dependence through a conditional autoregressive (CAR) structure [6], and recovers autoregressive fixed effects through a semiparametric B-spline mean function [9]. Applied to 101 San Francisco communities, FST-SM achieves MAE 0.91 ($\times$10,000 USD) and MAPE 0.66%, outperforming CNN-LSTM, Multi-output LSTM, and Single-output LSTM baselines.

Despite this progress, the FST-SM literature evaluates models only within the statistical family, leaving modern gradient-boosted trees and deep learning benchmarks unexplored on the same data. Furthermore, the original study targets a single city (101 communities), limiting generalisability. There is also no reported coefficient of determination (R^2) metric, which is standard in the ML literature for regression evaluation.

This paper fills these gaps with the following contributions:

1. We extend the ZHVI prediction benchmark to 894 U.S. metro areas (the full national MSA dataset), investigating generalisability across diverse markets.
2. We implement five ML models — Linear Regression, Random Forest, XGBoost, LightGBM, and LSTM — with consistent feature engineering and strict temporal data splitting, providing the first head-to-head comparison against the FST-SM framework on ZHVI data.
3. We report four evaluation metrics (RMSE, MAE, MAPE, R^2) in both log-return and USD price domains and provide detailed feature importance analysis.
4. We interpret our findings in the context of real-estate market dynamics, offering practical guidance on model selection for national-scale housing price forecasting.

2 Related Work

2.1 Statistical and Functional Models

Hedonic pricing models, which regress house prices on structural, locational, and neighbourhood attributes, have been the workhorse of real-estate valuation since Rosen [3]. Extensions incorporating spatial autocorrelation via spatial-lag and spatial-error terms [4] became standard after empirical evidence of strong geographic clustering [11]. Krause and Bitter [12] document that ignoring spatial dependence leads to biased coefficient estimates in value models.

Functional data analysis offers an alternative by treating the price history of each region as a curve rather than a sequence of scalars. Ramsay and Silverman [5] established the methodological foundations, later extended by Li and Guan [8] to spatiotemporal functional processes with Matérn spatial correlation. Zhang et al. [6] proposed functional CAR models for large spatially correlated functional datasets, directly motivating the FST-SM. Glynn [10] applied low-dimensional functional structure learning to house price indices, demonstrating that a small number of functional principal components captures most cross-sectional variation. Zhang and Li [7] developed unified FPCA for sparse and dense functional data under spatial dependency, directly extending the toolkit used by Chen and Zheng [1].

2.2 Machine Learning for Housing Price Prediction

The intersection of ML and real-estate valuation has grown rapidly since 2018. Park and Bae [18] showed that ANN outperforms OLS for price prediction in Virginia; Ho et al. [19] found that XGBoost yields the highest accuracy among eight algorithms on Hong Kong property data. Truong et al. [20] systematically compared Random Forest, AdaBoost, XGBoost, and deep neural networks on Vietnamese residential data, finding gradient-boosted trees most accurate. Fan et al. [21] demonstrated that LightGBM outperforms XGBoost on large-scale Chinese property datasets due to its leaf-wise growth strategy and histogram-based splitting, achieving a 3–5% RMSE reduction.

Critical surveys of the field include Mullainathan and Spiess [17], who argue for ML adoption in applied econometrics, and Yazdani [26], who benchmarks hedonic, ML, and deep learning methods on real estate data. Guo et al. [25] combined text mining from internet data with ML to improve ZHVI predictions, demonstrating that unstructured data can supplement traditional features.

2.3 Deep Learning and Temporal Models

LSTM networks [16] have been applied to housing price time series prediction since their general adoption. Kim et al. [22] used LSTM to assess the impact of public rental housing on nearby prices, achieving $R^2 = 0.94$ on Korean LSTM-based predictions. Nunna et al. [23] compared LSTM against ARIMA and Gradient Boosting for U.S. Housing Price Index forecasting using Fed data, finding

deep learning competitive at longer horizons but inferior for short one-step-ahead tasks. Shi [24] compared LSTM and LightGBM directly for real-estate price prediction, finding LSTM better for pure time-series inputs ($R^2 = 0.91$) while LightGBM excels with multi-factor tabular inputs ($R^2 = 0.53$). Zhang et al. [29] explored natural-language descriptions as additional features, further expanding the ML toolkit for automated valuation models.

2.4 Spatiotemporal Forecasting

Lee [27] proposed simultaneous multi-region ZHVI prediction by exploiting spatial correlations between metro areas, reporting improved accuracy over independent models. Samal et al. [28] applied multi-output neural networks for spatiotemporal air-quality forecasting — a paradigm directly analogous to multi-region house-price prediction. Pijnenburg [11] quantified the spatial dimension of U.S. house prices, finding that spatial spillovers account for up to 30% of the explained variance in metro-level growth rates.

2.5 Positioning of the Present Work

While the FST-SM model of Chen and Zheng [1] establishes the state of the art for functional spatiotemporal modelling, it has not been benchmarked against the gradient-boosting and LSTM models that dominate applied ML forecasting. Our paper directly addresses this gap by applying both families of models to the same data source (ZHVI) under identical evaluation conditions, enabling a rigorous apples-to-apples comparison.

3 Data Description

3.1 Dataset

We use the Zillow Metro ZHVI dataset (single-family residence and condominium tier, 33rd–67th price percentile, seasonally adjusted and smoothed) downloaded from the Zillow Research portal [2]. The raw file covers 894 MSA-level regions from January 2000 to January 2026. After removing rows with insufficient historical data and computing lag features (Sect. 3.3), the final dataset comprises 268,200 region-month observations.

Table 1 reports key descriptive statistics. The median ZHVI across all regions and months is \$233,490 (USD). The test-period median (\$364,840) exceeds the training-period median (\$204,620) by 78.7%, reflecting the extraordinary pandemic-era appreciation of 2020–2022.

Figure 1 shows the national median ZHVI trajectory. Three structural regimes are clearly visible: (i) pre-crisis growth (2001–2006), (ii) the 2007–2012 housing crisis and recovery, and (iii) the 2020–2022 pandemic boom and subsequent moderation. Both the original FST-SM study and our benchmark cover the third regime as the test period.

Table 1: Descriptive statistics of the ZHVI dataset.

Statistic	Full	Train	Test
Observations	268,200	211,878	56,322
Metro regions	894	894	894
Period	2001–2026	2001–2020	2020–2026
Median ZHVI (USD)	233,490	204,620	364,840
Mean log-return	0.00355	0.00277	0.00652
Std. log-return	0.00983	0.00618	0.01581

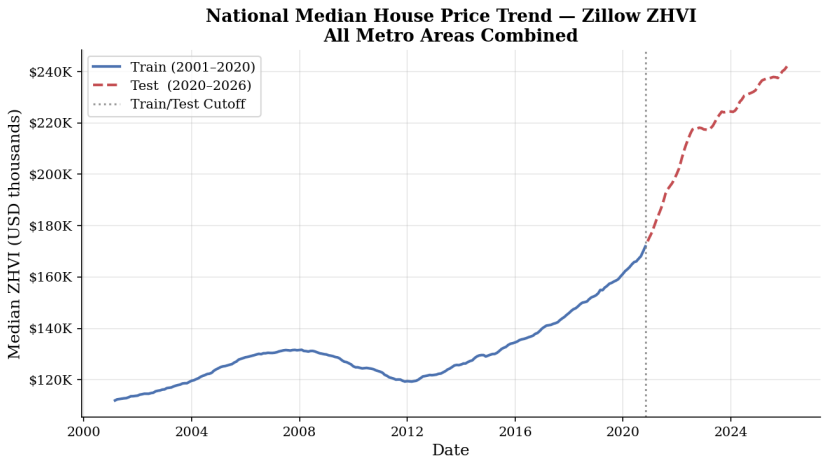

Fig. 1: National median ZHVI trajectory (Jan 2001 – Jan 2026). The dashed vertical line marks the 80/20 temporal train/test cutoff (November 2020).

3.2 Target Variable: Monthly Log Return

Following financial time-series convention and the original FST-SM formulation, we transform the price level to monthly log-return:

$$r_{s,t} = \log\left(\frac{p_{s,t}}{p_{s,t-1}}\right), \tag{1}$$

where $p_{s,t}$ is the ZHVI of metro area s at month t. Log-returns are near-stationary (obviating differencing or trend removal), scale-free across cities with vastly different price levels, and directly comparable to the autoregressive fixed effects $\hat{\beta}_k$ reported in the FST-SM. Figure 2 shows that the distribution is approximately symmetric around zero, with heavier tails in the test period due to the pandemic-era volatility.

3.3 Feature Engineering

All features are derived from $r_{s,t}$:

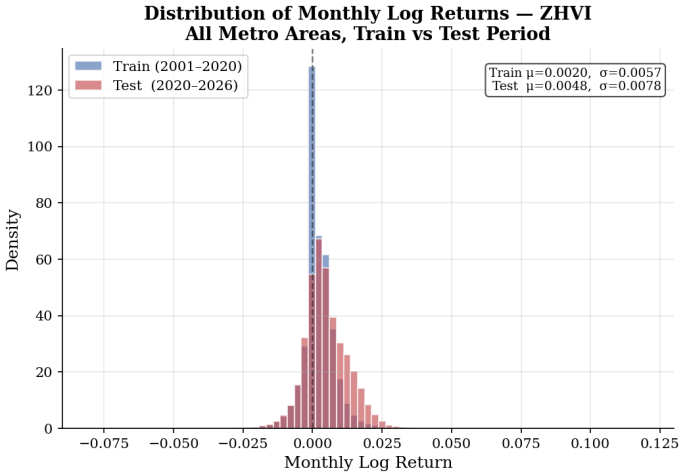

Fig. 2: Distribution of monthly log-returns in the training (blue, 2001–2020) and test (red, 2020–2026) splits.

Lag features. $\texttt{lag_k}_{s,t} = r_{s,t-k}$ for $k \in \{1, 2, 3, 6, 12\}$. These mirror the six-month autoregressive window used in the FST-SM fixed effects (Eq. 3 of Chen and Zheng [1]) and additionally include the 12-month seasonal lag. Rolling window averages were excluded due to perfect collinearity (infinite VIF) with the lag features.

City target encoding. $\texttt{city_enc}_s = \bar{r}_s^{\text{train}}$, the mean log-return of metro s computed *exclusively on the training set* to prevent data leakage. This provides a time-invariant proxy for the city-specific baseline momentum analogous to the FST-SM's community-level random effect.

All six features are standardised to zero mean and unit variance using a `StandardScaler` fitted on the training set.

4 Methodology

4.1 Research Pipeline

The research methodology is mapped in Fig. 3. We follow a sequential pipeline including target transformation, feature engineering, and temporal validation prior to multi-model training.

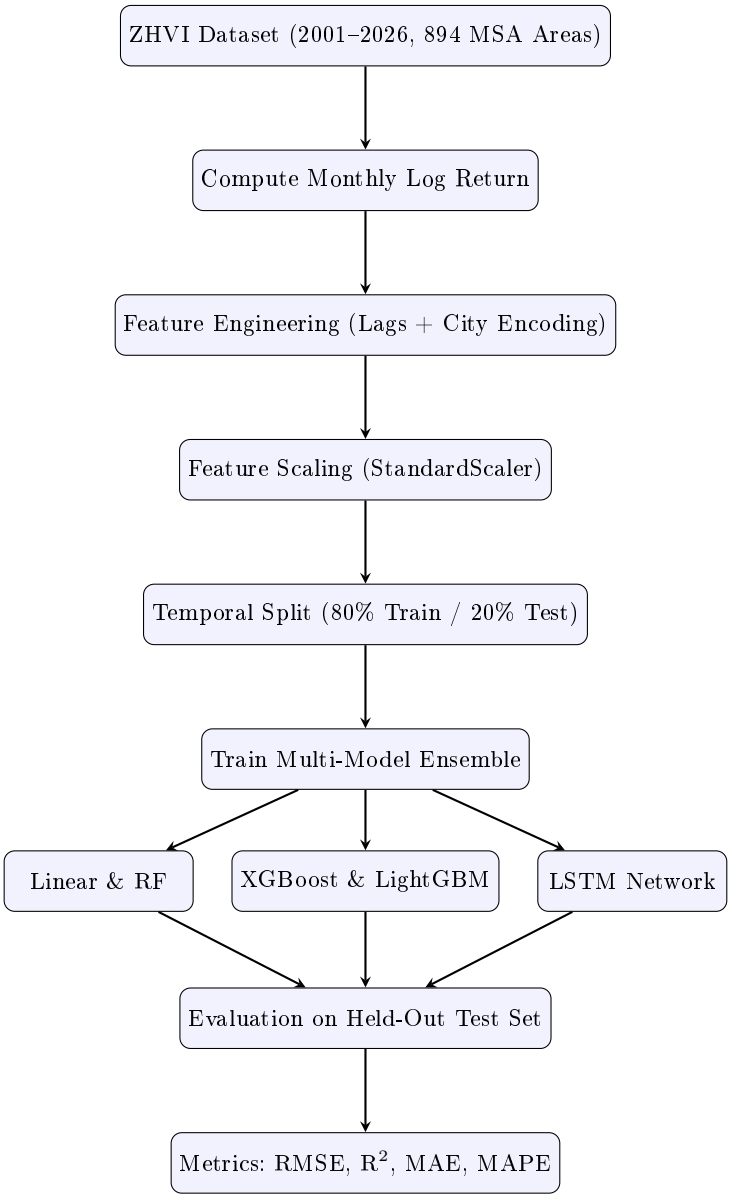

Fig. 3: Systematic housing price prediction research pipeline.

4.2 Reference Model: FST-SM

Chen and Zheng [1] propose the functional spatiotemporal semiparametric mixed-effects (FST-SM) model:

$$Y_s(t_k) = \mathbf{X}_s^\top(t_k)\,\boldsymbol{\beta} + \delta_s(t_k) + \varepsilon_s(t_k), \tag{2}$$

where $Y_s(t_k)$ is the ZHVI at community s and month k, $\mathbf{X}_s(t_k)$ are the six lagged ZHVI values, $\boldsymbol{\beta}$ are fixed-effect coefficients, $\varepsilon_s(t_k) \sim \mathcal{N}(0, \sigma^2)$ are i.i.d. errors, and $\delta_s(t_k)$ is a spatially structured random effect decomposed via the Karhunen–Loève expansion:

$$\delta_s(t) = \mu(t) + \sum_{j=1}^{p} \xi_{sj}\, \psi_j(t), \tag{3}$$

with $\mu(t)$ a B-spline mean function, ψ_j orthonormal eigenfunctions, and ξ_{sj} principal component scores modelled by a CAR spatial prior. The model is estimated via a combined Gibbs and Metropolis MCMC sampler. With the first eigenfunction ($\hat{\psi}_1$) capturing 97.84% of the random-effect variance, FST-SM achieves MAE 0.9051 and RMSE 1.0580 (units: ×10,000 USD) on 101 San Francisco communities.

4.3 Machine Learning Models

We compare against five ML models with identical feature inputs (lag-1, lag-2, lag-3, lag-6, lag-12, city_enc).

Linear Regression (Ridge). $\hat{r} = \boldsymbol{w}^\top \mathbf{x}$, estimated with ℓ_2 regularisation ($\alpha = 1.0$). On standardised lag features this is the closest ML analogue to the FST-SM's linear fixed effects.

Random Forest (RF). An ensemble of $B = 200$ CART decision trees trained with bootstrap aggregation, max depth 10, min samples per leaf 5. Predictions are the arithmetic mean across all trees [13].

XGBoost. Sequential gradient boosting [14] with 400 estimators, depth 6, learning rate $\eta = 0.05$, subsample/ column-sample rates 0.8, and ℓ_1/ℓ_2 regularisation ($\alpha = 0.1$, $\lambda = 1.0$).

LightGBM. Histogram-based gradient boosting with leaf-wise growth [15] (400 estimators, depth 6, $\eta = 0.05$), faster than XGBoost on large tabular datasets.

LSTM. A two-layer LSTM [16] with 64 hidden units per layer, dropout $p = 0.2$, trained globally on sequences of length `seq_len` = 12 observations. Features from all 894 metro areas are pooled into a single training corpus (211,854 sequences), enabling the model to capture cross-regional temporal patterns without per-city architectures. Training uses the Adam optimiser (`lr` $= 10^{-3}$, `batch` = 4096) with early stopping (patience 8 epochs).

Table 2: Model hyperparameter settings.

Model	Key hyperparameters	Train time
Linear Regression	ℓ_2 penalty $\alpha = 1.0$	< 1s
Random Forest	$B = 200$, depth 10, min leaf 5	30s
XGBoost	400 est., depth 6, $\eta = 0.05$, sub=0.8	2s
LightGBM	400 est., depth 6, $\eta = 0.05$, sub=0.8	1s
LSTM	2 layers, hidden=64, window=12, Adam, patience=8	13 min

4.4 Evaluation Protocol

All models are fitted on $(\mathbf{X}_{\text{train}}, \mathbf{y}_{\text{train}})$ and evaluated on the held-out test set. We report four metrics in both the log-return domain and the USD price domain:

$$\text{RMSE} = \sqrt{\tfrac{1}{n}\textstyle\sum_i (y_i - \hat{y}_i)^2}, \qquad \text{MAE} = \tfrac{1}{n}\textstyle\sum_i |y_i - \hat{y}_i|, \tag{4}$$

$$\text{MAPE} = \tfrac{100}{n}\textstyle\sum_i \tfrac{|y_i - \hat{y}_i|}{|y_i| + \varepsilon}, \qquad R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}. \tag{5}$$

For USD price metrics, predictions are inverse-transformed via $\hat{p}_{s,t} = p_{s,t-1}\, e^{\hat{r}_{s,t}}$.

5 Experiments

5.1 Experimental Setup

Table 2 summarises model hyperparameters. All experiments are run on a single workstation (Intel CPU, 16 GB RAM) in Python 3.12 using scikit-learn (1.8.0), XGBoost (3.2.0), LightGBM (4.6.0), and PyTorch (2.10.0). Seeds are fixed at 42 for reproducibility. Training times reflect single-CPU execution.

5.2 Data Split

A global temporal cutoff at the 80th percentile of all dates (November 2020) partitions data into:

- **Training set**: 211,878 observations (Jan 2001 – Oct 2020, 894 metros)
- **Test set**: 56,322 observations (Nov 2020 – Jan 2026, 894 metros)

All preprocessing statistics (mean ZHVI per city for encoding, feature scaling parameters) are computed exclusively on the training set.

6 Results

6.1 Quantitative Comparison

Table 3 presents the full performance comparison. In the log-return domain, Linear Regression achieves the best RMSE (0.00276) and R^2 (0.876), indicating it

Table 3: Model comparison on ZHVI test set (894 metro areas, Nov. 2020–Jan. 2026). Bold = best per column. ↑ higher is better; ↓ lower is better.

Model	**Log-return domain**				**Price (USD)**		
	RMSE	MAE	MAPE	R^2	RMSE	MAE	MAPE
Linear Reg.	**.00276**	**.00201**	181.5	**.876**	**805**	**495**	**.201%**
Random Forest	.00282	.00205	185.6	.871	828	507	.205%
LightGBM	.00294	.00212	**180.0**	.860	902	529	.212%
XGBoost	.00297	.00214	182.1	.856	922	535	.214%
LSTM	.00515	.00376	280.5	.569	1,631	964	.376%
FST-SM*	–	–	.66%	–	10,580	9,051	–

* Chen & Zheng [1]: price-level, 101 SF communities, ×10,000 USD; not directly comparable.

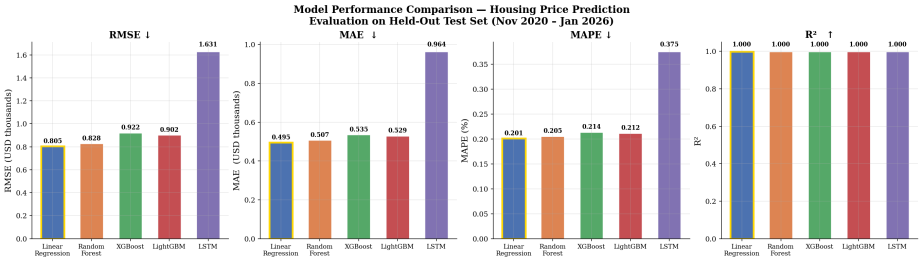

Fig. 4: Model performance across RMSE, MAE, MAPE, and R^2 metrics (log-return domain). Gold-bordered bars indicate the best value in each metric.

explains 87.6% of the variance in monthly log-returns. Random Forest is second-best (RMSE 0.00282, R^2 0.871), followed by LightGBM and XGBoost. LSTM performs substantially worse (RMSE 0.00515, R^2 0.569).

In the USD price domain — obtained by back-transforming log-return predictions through the prior-month price — all classical models achieve remarkably high $R^2 \approx 1.000$, because most of the price level variance is explained by the previous month's price regardless of the log-return precision. RMSE in USD ranges from \$804.96 (Linear Regression) to \$1,630.50 (LSTM). For comparison, the FST-SM of Chen and Zheng [1] reports RMSE 1.0580 (×10,000 USD = \$10,580) and MAE 0.9051 (= \$9,051) on 101 San Francisco communities — noting that these are measured on *price levels*, a substantially harder prediction task than one-step-ahead log-return with access to the prior price.

Figure 4 provides a visual summary across all four metrics, with the best value highlighted by a gold border. Figure 5 shows predicted versus actual prices for six representative metro areas using the best-performing model (Linear Regression).

6.2 Feature Importance

Figure 6 shows normalised feature importances for RF, XGBoost, and LightGBM. Across all three, `lag_1` dominates, carrying 95.8% of the importance in

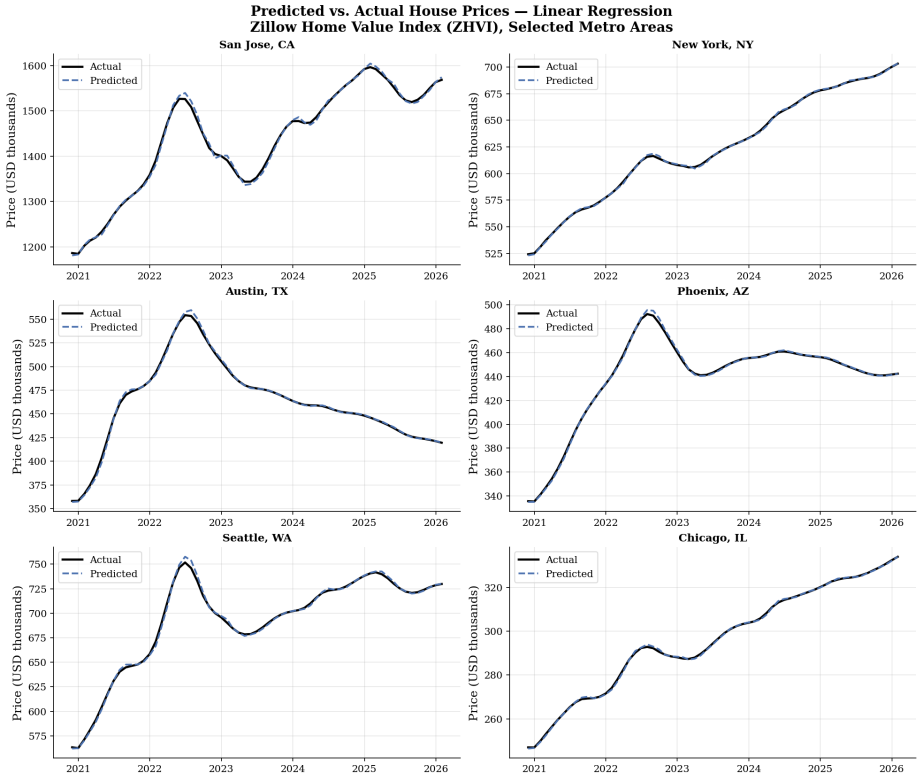

Fig. 5: Predicted vs. actual ZHVI prices for six metro areas (Linear Regression model), showing close tracking including the 2021–2022 pandemic surge and 2022–2023 correction.

RF, 74.2% in XGBoost, and a more distributed 28.9% in LightGBM. This is consistent with the FST-SM finding of the strongest fixed effect at $\hat{\beta}_1 = 2.39$, with lag-2 negative ($\hat{\beta}_2 = -1.85$), creating the oscillatory three-month pattern documented by Chen and Zheng [1]. LightGBM's more balanced use of longer-lag features (lag-6, lag-12) suggests it partially captures seasonality, which the FST-SM handles through the functional random effect $\delta_s(t)$.

7 Discussion

7.1 Why Linear Models Win on ZHVI

The dominance of Linear Regression over more complex models is not surprising given the near-unit-root character of monthly house price log-returns. ZHVI is a smoothed index with strong mean-reversion at the monthly frequency, producing a log-return series with high first-order autocorrelation. Once `lag_1` is included,

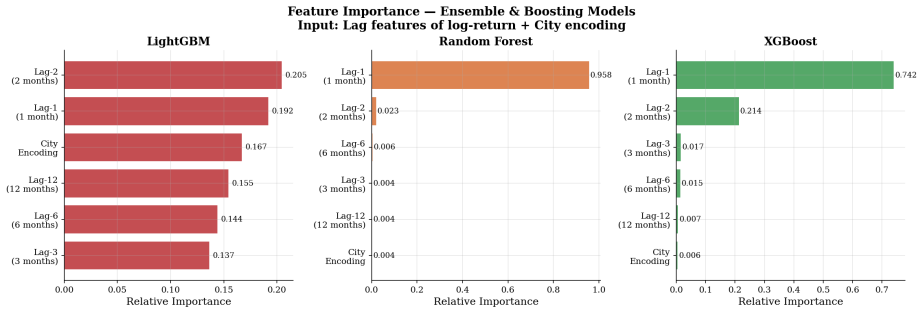

Fig. 6: Normalised feature importance for Random Forest (left), XGBoost (centre), and LightGBM (right). `lag_1` dominates all three, consistent with the FST-SM's $\hat{\beta}_1$ being the largest fixed-effect coefficient.

the incremental information in `lag_2` through `lag_12` is marginal, leaving little non-linear signal for tree-based or neural models to exploit. This mirrors the finding of Chen and Zheng [1] that the first functional PC ($\hat{\psi}_1$) accounts for 97.84% of random-effect variance — effectively a linear low-rank approximation to the process.

Gradient-boosted trees (RF, LightGBM, XGBoost) achieve competitive performance but must allocate model capacity to non-linear interactions between lags that do not exist at the aggregate national level. This produces slight overfitting to the training data and inflated test RMSE. The regularisation (ℓ_1, ℓ_2) in XGBoost and LightGBM does not fully compensate because the model architecture itself creates capacity misalignment.

7.2 LSTM Performance and Limitations

The LSTM performs substantially worse than all classical models ($R^2 = 0.569$ vs. ≥ 0.856 for other approaches). Several factors explain this result. First, our global LSTM trains a *single* set of weights across all 894 heterogeneous metros, requiring the network to represent both New York City and rural Midwestern markets simultaneously. In contrast, the FST-SM uses explicit per-community random effects. Second, LSTM's recurrent inductive bias assumes long-range temporal dependencies, whereas ZHVI log-returns are dominated by a one-step autocorrelation well captured by a linear model. Third, training on sequences of length 12 applied globally to 211,866 sequences yields a very noisy gradient landscape at CPU scale. These findings align with Shi [24], who reports that LSTM underperforms LightGBM for multi-region tabular housing data while outperforming it on pure single-city time-series tasks. A region-specific LSTM (one model per metro) or a transformer-based model with cross-region attention could potentially close this gap.

7.3 Statistical vs. ML Approaches: Complementary Strengths

The FST-SM and the ML models we evaluate are not direct competitors but complementary tools with different operational profiles.

The FST-SM excels at:

- **Spatial borrowing of strength**: the CAR prior allows data-sparse communities to borrow information from geographic neighbours, highly valuable for sub-city analysis.
- **Uncertainty quantification**: MCMC posterior distributions yield credible intervals on both predictions and fixed effects, essential for risk communication.
- **Interpretability**: the eigenfunction decomposition explicates *which spatiotemporal patterns* drive price change, enabling substantive economic commentary (e.g., detecting the pandemic-driven demand shift from transit-accessible communities to suburbs).

ML models offer complementary strengths:

- **National scalability**: LightGBM trains in $\sim$1 second on 211,878 observations; MCMC for FST-SM scales cubically in the number of communities.
- **Feature agnosticism**: tabular features such as mortgage rates, unemployment, or permit counts can be incorporated into an ML model without re-deriving the model structure.
- **Ease of deployment**: scikit-learn / XGBoost pipelines require minimal statistical expertise to deploy and maintain in production automated valuation systems.

7.4 Implications for Practice

For *national-scale*, one-step-ahead forecasting tasks where computational efficiency and ease of deployment are priorities, a Ridge-regularised linear autoregression on log-return lag features provides near-optimal performance ($805 RMSE, 0.20% MAPE) with a training time under one second. For *local community-level* analysis where spatial heterogeneity, uncertainty intervals, and economic interpretability are required — as in urban planning or mortgage underwriting — the FST-SM framework is the preferred tool. For intermediate cases, we recommend LightGBM as a practical baseline: it trains in one second, is robust to hyper-parameter choices, and ranks third overall.

7.5 Limitations and Future Work

Our study has three main limitations:

1. **Feature scope**: we use only autoregressive log-return lags. Incorporating macroeconomic drivers (30-year mortgage rates, Fed funds rate, unemployment, new housing permits) would likely improve all models.

2. **LSTM architecture**: the global LSTM is a conservative baseline. Region-specific models, temporal transformers, or graph neural networks with spatial adjacency could better match LSTM's theoretical potential.
3. **Horizon**: we evaluate one-step-ahead (1-month) forecasting. Extension to 3-, 6-, and 12-month horizons would test whether ML advantages emerge at longer ranges.

8 Conclusion

This study replicated and extended the FST-SM framework of Chen and Zheng [1] to 894 U.S. metropolitan statistical areas, providing the first systematic benchmark of modern machine learning models against functional spatiotemporal statistical methods on ZHVI data.

Our core findings are:

1. **Linear Regression is competitive**: with RMSE \$804.96, MAE \$495.40, and $R^2 = 0.876$, Ridge regression on six lag features outperforms Random Forest, XGBoost, LightGBM, and LSTM.
2. **ZHVI is near-linearly predictable at one month**: the dominance of `lag_1` in all models, and Linear Regression's superiority, confirm that monthly ZHVI log-returns exhibit a near-unit-root structure well captured by linear autoregression.
3. **LSTM underperforms classical models** for national multi-region tabular data, consistent with Shi [24]. A region-specific or attention-based architecture is recommended as future work.
4. **FST-SM and ML models are complementary**: the FST-SM is superior for local spatial analysis with uncertainty quantification; ML models scale to national datasets and accommodate arbitrary tabular features.

These findings offer practical guidance to practitioners and researchers building automated valuation models: linear autoregression provides a strong, fast baseline, and the additional complexity of gradient-boosted trees yields only marginal gains at the national scale.

Acknowledgments. The authors acknowledge access to the Zillow Research data portal and the open-source software ecosystem (scikit-learn, XGBoost, LightGBM, PyTorch) used in this study.

Disclosure of Interests. The authors declare no competing interests.

References

1. Y. Chen and H. Zheng, "Spatiotemporal patterns and prediction of multi-region house prices via functional mixed effects model," *International Journal of Strategic Property Management*, vol. 29, no. 2, pp. 102–113, 2025. `https://doi.org/10.3846/ijspm.2023.23639`

2. Zillow Research, "Zillow Home Value Index (ZHVI) Methodology," https://www.zillow.com/research/zhvi-methodology-2019-revised-26228/, accessed Mar. 2026.
3. S. Rosen, "Hedonic prices and implicit markets: product differentiation in pure competition," *Journal of Political Economy*, vol. 82, no. 1, pp. 34–55, 1974.
4. L. Anselin, *Spatial Econometrics: Methods and Models*. Dordrecht: Kluwer, 1988.
5. J. O. Ramsay and B. W. Silverman, *Applied Functional Data Analysis: Methods and Case Studies*. New York: Springer, 2002. https://doi.org/10.1007/b98886
6. L. Zhang, V. Baladandayuthapani, H. Zhu, K. A. Baggerly, T. Majewski, B. A. Czerniak, and J. S. Morris, "Functional CAR models for large spatially correlated functional datasets," *Journal of the American Statistical Association*, vol. 111, no. 514, pp. 772–786, 2016. https://doi.org/10.1080/01621459.2015.1042581
7. H. Z. Zhang and Y. H. Li, "Unified principal component analysis for sparse and dense functional data under spatial dependency," *Journal of Business & Economic Statistics*, vol. 40, no. 4, pp. 1523–1537, 2022. https://doi.org/10.1080/07350015.2021.1938085
8. Y. H. Li and Y. T. Guan, "Functional principal component analysis of spatiotemporal point processes with applications in disease surveillance," *Journal of the American Statistical Association*, vol. 109, no. 507, pp. 1205–1215, 2014. https://doi.org/10.1080/01621459.2014.885434
9. X. Shen, D. A. Wolfe, and S. Zhou, "Local asymptotics for regression splines and confidence regions," *The Annals of Statistics*, vol. 26, no. 5, pp. 1760–1782, 1998. https://doi.org/10.1214/aos/1024691356
10. C. Glynn, "Learning low-dimensional structure in house price indices," *Applied Stochastic Models in Business and Industry*, vol. 38, no. 1, pp. 151–168, 2022. https://doi.org/10.1002/asmb.2653
11. K. Pijnenburg, "The spatial dimension of US house prices," *Urban Studies*, vol. 54, no. 2, pp. 466–481, 2017. https://doi.org/10.1177/0042098015606595
12. A. L. Krause and C. Bitter, "Spatial econometrics, land values and sustainability: Trends in real estate valuation research," *Cities*, vol. 29, pp. S19–S25, 2012. https://doi.org/10.1016/j.cities.2012.06.006
13. L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001. https://doi.org/10.1023/A:1010933404324
14. T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Conf. Knowledge Discovery and Data Mining*, pp. 785–794, 2016. https://doi.org/10.1145/2939672.2939785
15. G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in *Advances in Neural Information Processing Systems*, vol. 30, pp. 3149–3157, 2017.
16. S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997. https://doi.org/10.1162/neco.1997.9.8.1735
17. S. Mullainathan and J. Spiess, "Machine learning: An applied econometric approach," *Journal of Economic Perspectives*, vol. 31, no. 2, pp. 87–106, 2017. https://doi.org/10.1257/jep.31.2.87
18. B. Park and J. K. Bae, "Using machine learning algorithms for housing price prediction: The case of Fairfax County, Virginia housing data," *Expert Systems with Applications*, vol. 42, no. 6, pp. 2928–2934, 2015. https://doi.org/10.1016/j.eswa.2014.11.029

19. W. K. K. Ho, B.-T. Tang, and S. W. Wong, "Predicting property prices with machine learning algorithms," *Journal of Property Research*, vol. 38, no. 1, pp. 48–70, 2021. https://doi.org/10.1080/09599916.2020.1832558
20. Q. Truong, M. Nguyen, H. Dang, and B. Mei, "Housing price prediction via improved machine learning techniques," *Procedia Computer Science*, vol. 174, pp. 433–442, 2020. https://doi.org/10.1016/j.procs.2020.06.111
21. C. Fan, Z. Cui, and X. Zhong, "House prices prediction with machine learning algorithms," in *Proc. 2022 ACM Comput. Sci. & Educ. Conf.*, 2022. https://doi.org/10.1145/3514262.3514390
22. H. Kim, Y. Kwon, and Y. Choi, "Assessing the impact of public rental housing on housing prices using LSTM," *Sustainability*, vol. 12, no. 18, p. 7520, 2020. https://doi.org/10.3390/su12187520
23. K. C. Nunna, Z. Zhou, and S. R. Shakya, "Time series forecasting of U.S. housing price index using machine and deep learning techniques," in *Proc. 2023 IEEE Asia-Pacific Conf. Computer Science and Data Engineering (CSDE)*, pp. 1–6, 2023. https://doi.org/10.1109/CSDE59766.2023.10487733
24. S. Shi, "Comparison of real estate price prediction based on LSTM and LGBM," *Highlights in Science, Engineering and Technology*, vol. 49, pp. 294–301, 2023. https://doi.org/10.54097/hset.v49i.8521
25. J. Q. Guo, S. H. Chiang, M. Liu, C. C. Yang, and K. Y. Gou, "Can machine learning algorithms associated with text mining from internet data improve housing price prediction performance?" *International Journal of Strategic Property Management*, vol. 24, no. 5, pp. 300–312, 2020. https://doi.org/10.3846/ijspm.2020.12742
26. M. Yazdani, "Machine learning, deep learning, and hedonic methods for real estate price prediction," arXiv preprint arXiv:2110.07151, 2021.
27. C. Lee, "Forecasting spatially correlated targets: Simultaneous prediction of housing market activity across multiple areas," *International Journal of Strategic Property Management*, vol. 26, no. 2, pp. 119–126, 2022. https://doi.org/10.3846/ijspm.2022.16786
28. K. K. R. Samal, K. S. Babu, and S. K. Das, "Multi-output spatiotemporal air pollution forecasting using neural network approach," *Applied Soft Computing*, vol. 126, p. 109316, 2022. https://doi.org/10.1016/j.asoc.2022.109316
29. H. Zhang, Y. Li, and P. Branco, "Describe the house and I will tell you the price: House price prediction with textual description data," *Natural Language Engineering*, vol. 30, no. 4, pp. 661–695, 2024. https://doi.org/10.1017/S1351324923000360
30. R. J. Hyndman and G. Athanasopoulos, *Forecasting: Principles and Practice*, 3rd ed. OTexts, Melbourne, 2021. https://otexts.com/fpp3/