

Functional Semiparametric and Machine Learning Ensembles for Large Scale Predictive High Granularity Forecasting

Nguyen Trinh Quy^[0009-0004-8203-0736], Thu Le^[0009-0008-3480-8311],
Nguyen Cong Minh^[0009-0009-6348-0359], Nguyen Do Nhat
Quang^[0009-0000-0055-3957], and Le Pham Hue Chi^[0009-0002-7550-9234]

FPT University HCM, Ho Chi Minh City, Vietnam
nguyentrinhquy1411@gmail.com, thulvm@fpt.edu.vn, {nguyenminhfptu,
nguyendonhatquangaidev, lephamhuechi}@gmail.com

Abstract. This paper systematically compares the functional spatiotemporal semiparametric mixed-effects (FST-SM) model [1] against five machine learning (ML) approaches for forecasting multi-region house prices on the Zillow Home Value Index (ZHVI). Analyzing 894 U.S. metropolitan areas over a 25-year period, we find that robust linear autoregression achieves the lowest log-return RMSE (0.00276) and highest R^2 (0.876), outperforming tree-based and recurrent neural network models at the national scale due to the near-unit-root temporal structure. Our central contribution is providing a rigorous benchmark that delineates the practical complementarity of these methods: statistical modeling excels at localized spatial uncertainty quantification, while efficient scalable machine learning pipelines are competitive for broad aggregate forecasting.

Keywords: Housing price prediction, Zillow ZHVI, Functional mixed-effects model, Machine learning, Gradient boosting, LSTM, Spatiotemporal analysis

1 Introduction

Residential real-estate markets exhibit pronounced spatiotemporal structure. Prices in geographically proximate markets tend to co-move, while long-run trajectories diverge substantially across cities and communities. Accurate, scalable prediction of these dynamics is economically significant: it informs mortgage risk assessment, zoning decisions, urban investment, and household financial planning [12].

The Zillow Home Value Index (ZHVI) is a canonical benchmark for this problem. It provides monthly, seasonally adjusted, smoothed median home values at the metro-statistical area (MSA) level for more than 900 U.S. cities, covering over two decades of market fluctuations including the 2007 to 2012 financial

crisis, the COVID-19 pandemic price surge of 2020 to 2022, and the subsequent correction [2].

Statistical approaches to spatiotemporal housing price modelling have a rich history. Hedonic regression [3] accounts for property attribute heterogeneity; spatial-lag models [4] capture geographic autocorrelation; and functional data analysis (FDA) methods [5] represent price trajectories as continuous curves. The recent FST-SM model of Chen and Zheng [1] elegantly unifies these threads: it decomposes ZHVI trajectories via functional principal component analysis (FPCA), models spatial dependence through a conditional autoregressive (CAR) structure [6], and recovers autoregressive fixed effects through a semiparametric B-spline mean function [9]. Applied to 101 San Francisco communities, FST-SM achieves MAE 0.91 ($\times 10,000$ USD) and MAPE 0.66%, outperforming CNN-LSTM, Multi-output LSTM, and Single-output LSTM baselines.

Despite this progress, the FST-SM literature evaluates models only within the statistical family, leaving modern gradient-boosted trees and deep learning benchmarks unexplored on the same data. Furthermore, the original study targets a single city (101 communities), limiting generalisability. There is also no reported coefficient of determination (R^2) metric, which is standard in the ML literature for regression evaluation.

This paper fills these gaps by establishing several main contributions. First, we extend the ZHVI prediction benchmark to 894 U.S. metro areas to investigate generalisability across diverse markets. Second, we implement five ML models (Linear Regression, Random Forest, XGBoost, LightGBM, and LSTM) using consistent feature engineering and strict temporal data splitting, enabling the first head-to-head comparison against the FST-SM framework on ZHVI data. Third, we report comprehensive evaluation metrics across log-return and USD price domains alongside detailed feature importance analysis. Finally, we interpret our findings in the context of real-estate market dynamics, offering practical guidance on model selection for national-scale housing price forecasting.

2 Related Work

Statistical approaches to real-estate valuation have rapidly evolved from foundational hedonic regression estimating structural attributes [3] to highly advanced models adept at capturing profound spatial and temporal dependencies [4,11]. Furthermore, functional data analysis portrays price histories continuously [5], with notable innovations marrying functional principal component analysis and conditional autoregressive structures to properly accommodate large, spatially correlated datasets [6,8]. The functional spatiotemporal semiparametric mixed-effects (FST-SM) model [1] has specifically integrated many of these threads to successfully reflect detailed local dynamics. Conversely, expansive machine learning architectures, particularly gradient-boosted methods (like XGBoost and LightGBM) and specialized deep neural networks (e.g., LSTMs), increasingly dominate predictive benchmarks due to their capacity to absorb highly non-linear feature sets across massive data pools efficiently [19,21,24]. Thus, robust

multi-region spatiotemporal tasks are beginning to leverage key successes drawn from both analytical frameworks [27,28].

Our research purposefully bridges the structural gap between sophisticated statistical procedures and computationally scalable ML methods. Rather than separately assessing approaches within their individual domains as commonly observed in the prior literature, we establish a rigorous comparative benchmark setting the FST-SM against five major computational baselines on unified national constraints. By comprehensively standardizing feature inputs alongside temporally isolated parameters across an extensive 894-metro sample, our project clarifies critical trade-offs between essential linear autoregression, non-linear algorithmic ensembles, and expansive deep learning systems. Above all, this investigation directly contributes clarity to the broader literature by illuminating where complex spatial biases distinctly outvalue broader computational structures versus when simpler algorithmic pipelines operate functionally optimally for scalable national economic forecasting.

3 Proposed Methodology

3.1 Research Pipeline

The research methodology is mapped in Fig. 1. We follow a sequential pipeline including target transformation, feature engineering, and temporal validation prior to multi-model training.




Fig. 1: Comparative research methodology with parallel statistical and ML paths.

3.2 Reference Model: FST-SM

Chen and Zheng [1] propose the functional spatiotemporal semiparametric mixed-effects (FST-SM) model:

$$Y_s(t_k) = \mathbf{X}_s^\top(t_k) \boldsymbol{\beta} + \delta_s(t_k) + \varepsilon_s(t_k), \quad (1)$$

where $Y_s(t_k)$ is the ZHVI at community s and month k , $\mathbf{X}_s(t_k)$ are the six lagged ZHVI values, $\boldsymbol{\beta}$ are fixed-effect coefficients, $\varepsilon_s(t_k) \sim \mathcal{N}(0, \sigma^2)$ are i.i.d. errors, and $\delta_s(t_k)$ is a spatially structured random effect decomposed via the Karhunen to Loève expansion:

$$\delta_s(t) = \mu(t) + \sum_{j=1}^p \xi_{sj} \psi_j(t), \quad (2)$$

with $\mu(t)$ a B-spline mean function, ψ_j orthonormal eigenfunctions, and ξ_{sj} principal component scores modelled by a CAR spatial prior. The model is estimated via a combined Gibbs and Metropolis MCMC sampler. With the first eigenfunction ($\hat{\psi}_1$) capturing 97.84% of the random-effect variance, FST-SM achieves MAE 0.9051 and RMSE 1.0580 (units: $\times 10,000$ USD) on 101 San Francisco communities.

3.3 Machine Learning Models

We compare against five ML models with identical feature inputs (lag-1, lag-2, lag-3, lag-6, lag-12, city_enc).

Linear Regression (Ridge). $\hat{r} = \mathbf{w}^\top \mathbf{x}$, estimated with ℓ_2 regularisation ($\alpha = 1.0$). On standardised lag features this is the closest ML analogue to the FST-SM’s linear fixed effects.

Random Forest (RF). An ensemble of $B = 200$ CART decision trees trained with bootstrap aggregation, max depth 10, min samples per leaf 5. Predictions are the arithmetic mean across all trees [13].

XGBoost. Sequential gradient boosting [14] with 400 estimators, depth 6, learning rate $\eta = 0.05$, subsample/ column-sample rates 0.8, and ℓ_1/ℓ_2 regularisation ($\alpha = 0.1$, $\lambda = 1.0$).

LightGBM. Histogram-based gradient boosting with leaf-wise growth [15] (400 estimators, depth 6, $\eta = 0.05$), faster than XGBoost on large tabular datasets.

LSTM. A two-layer LSTM [16] with 64 hidden units per layer, dropout $p = 0.2$, trained globally on sequences of length `seq_len` = 12 observations. Features from all 894 metro areas are pooled into a single training corpus (211,854 sequences), enabling the model to capture cross-regional temporal patterns without per-city architectures. Training uses the Adam optimiser (`lr` = 10^{-3} , `batch` = 4096) with early stopping (patience 8 epochs).

4 Experimental Setup

4.1 Dataset Description and Feature Engineering

We use the Zillow Metro ZHVI dataset (single-family residence and condominium tier, 33rd–67th price percentile, seasonally adjusted and smoothed) downloaded from the Zillow Research portal [2]. The raw file covers 894 MSA-level regions from January 2000 to January 2026. After removing rows with insufficient historical data and computing lag features, the final dataset comprises 268,200 region-month observations.

Table 1 reports key descriptive statistics. The median ZHVI across all regions and months is \$233,490 (USD). The test-period median (\$364,840) exceeds the training-period median (\$204,620) by 78.7%, reflecting the extraordinary pandemic-era appreciation of 2020 to 2022.

Table 1: Descriptive statistics of the ZHVI dataset across full, training, and test periods.

Statistic	Full	Train	Test
Observations	268,200	211,878	56,322
Metro regions	894	894	894
Period	2001–2026	2001–2020	2020–2026
Median ZHVI (USD)	233,490	204,620	364,840
Mean log-return	0.00355	0.00277	0.00652
Std. log-return	0.00983	0.00618	0.01581

Figure 2 shows the national median ZHVI trajectory, while Fig. 3 illustrates the distribution of log-returns exhibiting characteristic heavy tails and high kurtosis.




Fig. 2: National median ZHVI trajectory.




Fig. 3: ZHVI log-return distribution.

Following financial time-series convention and the original FST-SM formulation, we transform the price level to monthly log-return $r_{s,t} = \log(p_{s,t}/p_{s,t-1})$.

Table 2: Model hyperparameters and approximate single-machine training times.

Model	Key hyperparameters	Train time
Linear Regression	ℓ_2 penalty $\alpha = 1.0$	< 1s
Random Forest	$B = 200$, depth 10, min leaf 5	30s
XGBoost	400 est., depth 6, $\eta = 0.05$, sub=0.8	2s
LightGBM	400 est., depth 6, $\eta = 0.05$, sub=0.8	1s
LSTM	2 layers, hidden=64, window=12, Adam, patience=8	13 min

All features are derived from $r_{s,t}$, including lag features (`lag_ks,t` for $k \in \{1, 2, 3, 6, 12\}$) and city target encoding (`city_encs`). All six features are standardised to zero mean and unit variance using a `StandardScaler` fitted on the training set.

4.2 Implementation Details and Data Split

Table 2 summarises model hyperparameters. All experiments are run on a single workstation (Intel CPU, 16 GB RAM) in Python 3.12 using scikit-learn (1.8.0), XGBoost (3.2.0), LightGBM (4.6.0), and PyTorch (2.10.0).

A global temporal cutoff at the 80th percentile of all dates (November 2020) partitions the data into a training set comprising 211,878 separate observations alongside a withheld evaluation test set storing 56,322 final entries.

4.3 Evaluation Metrics

We report four metrics in both the log-return domain and the USD price domain:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_i (y_i - \hat{y}_i)^2}, \quad \text{MAE} = \frac{1}{n} \sum_i |y_i - \hat{y}_i|, \quad (3)$$

$$\text{MAPE} = \frac{100}{n} \sum_i \frac{|y_i - \hat{y}_i|}{|y_i| + \varepsilon}, \quad R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}. \quad (4)$$

For USD price metrics, predictions are inverse-transformed via $\hat{p}_{s,t} = p_{s,t-1} e^{\hat{r}_{s,t}}$.

5 Results and Discussion

5.1 Quantitative Performance Analysis

Table 3 presents the full performance comparison. In the log-return domain, Linear Regression achieves the best RMSE (0.00276) and R^2 (0.876), indicating it explains 87.6% of the variance in monthly log-returns. Random Forest is second-best (RMSE 0.00282, R^2 0.871), followed by LightGBM and XGBoost. LSTM performs substantially worse (RMSE 0.00515, R^2 0.569). Figure 4 visualises the cross-model ranking on all three metrics.

Table 3: One-step-ahead forecast performance comparison across all models. Bold denotes the best result per metric. FST-SM results are price-level only and not directly comparable (see footnote).

Model	Log-return domain				Price (USD)		
	RMSE	MAE	MAPE	R ²	RMSE	MAE	MAPE
Linear Reg.	.00276	.00201	181.5	.876	805	495	.201%
Random Forest	.00282	.00205	185.6	.871	828	507	.205%
LightGBM	.00294	.00212	180.0	.860	902	529	.212%
XGBoost	.00297	.00214	182.1	.856	922	535	.214%
LSTM	.00515	.00376	280.5	.569	1,631	964	.376%
FST-SM*	to	to	.66%	to	10,580	9,051	to

* Chen & Zheng [1]: price-level MAE 0.9051, RMSE 1.0580 ($\times 10,000$ USD), 101 SF communities; not directly comparable to our national log-return evaluation.

In the USD price domain, obtained by back-transforming log-return predictions through the prior-month price, all classical models achieve remarkably high $R^2 \approx 1.000$, because most of the price level variance is explained by the previous month’s price regardless of the log-return precision. RMSE in USD ranges from \$804.96 (Linear Regression) to \$1,630.50 (LSTM). Figure 5 illustrates predicted versus actual prices for representative metro areas. For comparison, the FST-SM of Chen and Zheng [1] reports RMSE 1.0580 ($\times 10,000$ USD = \$10,580) and MAE 0.9051 ($\times 10,000$ USD = \$9,051), measured on price levels for 101 San Francisco communities and thus not directly comparable to our national log-return benchmark.




Fig. 4: Model performance comparison (RMSE, MAE, R²).

5.2 Feature Importance and Interpretability

Figure 6 shows normalised feature importances for RF, XGBoost, and LightGBM. Across all three, `lag_1` dominates, carrying 95.8% of the importance in RF, 74.2% in XGBoost, and a more distributed 28.9% in LightGBM. This is consistent with the FST-SM finding of the strongest fixed effect at $\hat{\beta}_1 = 2.39$, with lag-2 negative ($\hat{\beta}_2 = -1.85$), creating the oscillatory three-month pattern




Fig. 5: Predicted vs. actual ZHVI prices for representative metro areas.

documented by Chen and Zheng [1]. LightGBM's more balanced use of longer-lag features (lag-6, lag-12) suggests it partially captures seasonality, which the FST-SM handles through the functional random effect $\delta_s(t)$.




Fig. 6: Normalised feature importance.

5.3 Discussion and Implications

The dominance of Linear Regression over more complex models is not surprising given the near-unit-root character of monthly house price log-returns. ZHVI is a

smoothed index with strong mean-reversion at the monthly frequency, producing a log-return series with high first-order autocorrelation. Once `lag_1` is included, the incremental information in `lag_2` through `lag_12` is marginal, leaving little non-linear signal for tree-based or neural models to exploit. This mirrors the finding of Chen and Zheng [1] that the first functional PC ($\hat{\psi}_1$) accounts for 97.84% of random-effect variance, effectively a linear low-rank approximation to the process.

Gradient-boosted trees (RF, LightGBM, XGBoost) achieve competitive performance but must allocate model capacity to non-linear interactions between lags that do not exist at the aggregate national level. This produces slight overfitting to the training data and inflated test RMSE. The regularisation (ℓ_1 , ℓ_2) in XGBoost and LightGBM does not fully compensate because the model architecture itself creates capacity misalignment.

The LSTM performs substantially worse than all classical models ($R^2 = 0.569$ vs. ≥ 0.856 for other approaches). Several factors explain this result. First, our global LSTM trains a *single* set of weights across all 894 heterogeneous metros, requiring the network to represent both New York City and rural Midwestern markets simultaneously. In contrast, the FST-SM uses explicit per-community random effects. Second, LSTM’s recurrent inductive bias assumes long-range temporal dependencies, whereas ZHVI log-returns are dominated by a one-step autocorrelation well captured by a linear model. Third, training on sequences of length 12 applied globally to 211,866 sequences yields a very noisy gradient landscape at CPU scale. These findings align with Shi [24], who reports that LSTM underperforms LightGBM for multi-region tabular housing data while outperforming it on pure single-city time-series tasks. A region-specific LSTM (one model per metro) or a transformer-based model with cross-region attention could potentially close this gap.

The FST-SM and the ML models we evaluate are not direct competitors but complementary tools with different operational profiles.

The FST-SM excels at three tasks: (i) *spatial borrowing of strength*, where the CAR prior lets data-sparse communities leverage information from geographic neighbours; (ii) *uncertainty quantification*, as MCMC posterior distributions yield credible intervals on predictions and fixed effects; and (iii) *interpretability*, since eigenfunction decompositions reveal the dominant spatiotemporal patterns. ML models are complementary along different axes: gradient-boosted pipelines (LightGBM, XGBoost) train on hundreds of thousands of observations in seconds and ingest additional tabular inputs—mortgage rates, employment figures—without architectural changes, while accessible frameworks (scikit-learn, PyTorch) lower the barrier to production deployment.

In conclusion, while the FST-SM remains an excellent theoretical baseline, the practical application of Gradient-Boosted ensembles and sequential LSTM networks offers superior real-time predictive value in high volatility housing markets.

To support reproducible research and facilitate further benchmarking, the complete implementation, datasets, and preprocessing scripts are publicly available at <https://github.com/nguyentrinhquy1411/ady-spring2026>.

5.4 Multi-Step-Ahead Forecasting

Following the reviewer recommendation and our single-step baseline, we extend the evaluation to 2-, 6-, and 12-month direct-forecast horizons. For each horizon $h \in \{2, 6, 12\}$, the target is $r_{s,t+h} = \log(p_{s,t+h}/p_{s,t+h-1})$ and the same six lag features (computed at time t) serve as inputs. Models are retrained from scratch under the identical temporal split; no future information leakage occurs. Table 4 reports log-return RMSE and R^2 averaged across all 894 metro areas.

Table 4: Multi-step-ahead log-return RMSE and R^2 for $h \in \{2, 6, 12\}$ months. Bold denotes the best result per column.

Model	$h = 2$		$h = 6$		$h = 12$	
	RMSE	R^2	RMSE	R^2	RMSE	R^2
Linear Reg.	.00431	.614	.00892	.358	.01104	.147
Random Forest	.00449	.583	.00871	.388	.01083	.176
LightGBM	.00441	.597	.00858	.406	.01061	.208
XGBoost	.00453	.577	.00864	.398	.01069	.197
LSTM	.00618	.382	.01053	.211	.01287	.021

Several patterns emerge. First, forecast accuracy degrades sharply as the horizon lengthens: R^2 drops from approximately 0.60 at $h = 2$ to below 0.21 at $h = 12$ for all models, confirming that ZHVI log-returns are largely unpredictable beyond a few months. Second, the ranking flips at longer horizons: while Linear Regression dominates at $h = 2$, LightGBM becomes the best model at $h = 6$ and $h = 12$, indicating that non-linear lag interactions gain marginal relevance as the autocorrelation signal from `lag_1` decays. Third, LSTM continues to lag, with near-zero R^2 at $h = 12$, consistent with its difficulty in handling the sparse long-range signal in a global pooling regime. These results suggest that practitioners requiring medium-to-long-horizon forecasts should prefer gradient-boosted ensembles over both linear and recurrent architectures.

5.5 Limitations

Our study has three remaining limitations that guide future work. Feature scope is restricted to autoregressive log-return lags; incorporating macroeconomic co-

variates (unemployment, 30-year mortgage rates, housing permits) should improve all models. The global LSTM represents a conservative deep-learning baseline; region-specific LSTMs, temporal transformers, or spatial graph networks may recover additional predictive signal. Finally, the direct multi-step evaluation above assumes fixed lag features computed at time t ; iterative and sequence-to-sequence strategies could provide complementary insights.

6 Conclusion

This study significantly replicated and extended the FST-SM framework by evaluating it robustly alongside five diverse machine learning solutions across 894 metropolitan statistical zones. We successfully demonstrated that simplistic linear regression maintains exceptionally competitive benchmarks, securing high R^2 values and consistently outstripping complicated Random Forest, LightGBM, and LSTM counterparts during rigorous trials. These robust metrics largely confirm that single-month ZHVI parameters function predictably around proximate linear unit constraints predominantly resolved through lag indicators. Consequently, universal deep architectures like LSTM perform less optimally when universally modeling broad independent temporal tables without granular community parameters. Taken together, our comparative findings underscore the fundamentally complementary nature uniting these disciplines. While comprehensive statistical techniques excel significantly at resolving specific localized uncertainties structurally, massive modern ML counterparts offer rapid frictionless deployment over aggregate national contexts. Collectively, adopting straightforward computational linear setups represents a vastly powerful optimization mechanism when designing automated real-estate forecasting applications, showing that leveraging extensive deep networks may generate distinctly nominal utility during massive structural market applications.

Acknowledgments. The authors acknowledge access to the Zillow Research data portal and the open-source software ecosystem (scikit-learn, XGBoost, LightGBM, PyTorch) used in this study.

Disclosure of Interests. The authors declare no competing interests.

References

1. Y. Chen and H. Zheng, "Spatiotemporal patterns and prediction of multi-region house prices via functional mixed effects model," *International Journal of Strategic Property Management*, vol. 29, no. 2, pp. 102 to 113, 2025. Available: <https://journals.vilniustech.lt/index.php/IJSPM/article/view/23639>
2. Zillow Research, "Zillow Home Value Index (ZHVI) Methodology," accessed Mar. 2026.
3. S. Rosen, "Hedonic prices and implicit markets: product differentiation in pure competition," *Journal of Political Economy*, vol. 82, no. 1, pp. 34 to 55, 1974.
4. L. Anselin, *Spatial Econometrics: Methods and Models*. Dordrecht: Kluwer, 1988.
5. J. O. Ramsay and B. W. Silverman, *Applied Functional Data Analysis: Methods and Case Studies*. New York: Springer, 2002. <https://doi.org/10.1007/b98886>
6. L. Zhang, V. Baladandayuthapani, H. Zhu, K. A. Baggerly, T. Majewski, B. A. Czerniak, and J. S. Morris, "Functional CAR models for large spatially correlated functional datasets," *Journal of the American Statistical Association*, vol. 111, no. 514, pp. 772 to 786, 2016. <https://doi.org/10.1080/01621459.2015.1042581>
7. H. Z. Zhang and Y. H. Li, "Unified principal component analysis for sparse and dense functional data under spatial dependency," *Journal of Business & Economic Statistics*, vol. 40, no. 4, pp. 1523 to 1537, 2022. <https://doi.org/10.1080/07350015.2021.1938085>
8. Y. H. Li and Y. T. Guan, "Functional principal component analysis of spatiotemporal point processes with applications in disease surveillance," *Journal of the American Statistical Association*, vol. 109, no. 507, pp. 1205 to 1215, 2014. <https://doi.org/10.1080/01621459.2014.885434>
9. X. Shen, D. A. Wolfe, and S. Zhou, "Local asymptotics for regression splines and confidence regions," *The Annals of Statistics*, vol. 26, no. 5, pp. 1760 to 1782, 1998.
10. C. Glynn, "Learning low-dimensional structure in house price indices," *Applied Stochastic Models in Business and Industry*, vol. 38, no. 1, pp. 151 to 168, 2022.
11. K. Pijnenburg, "The spatial dimension of US house prices," *Urban Studies*, 2017.
12. A. L. Krause and C. Bitter, "Spatial econometrics, land values and sustainability," *Cities*, 2012.
13. L. Breiman, "Random forests," *Machine Learning*, 2001.
14. T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Conf. KDD*, 2016.
15. G. Ke et al., "LightGBM: A highly efficient gradient boosting tree," in *NeurIPS*, 2017.
16. S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, 1997.
17. S. Mullainathan and J. Spiess, "Machine learning: An applied econometric approach," *JEP*, 2017.
18. B. Park and J. Bae, "Using ML for housing price prediction," *ESWA*, 2015.
19. W. Ho et al., "Predicting property prices with ML," *JPR*, 2021.
20. Q. Truong et al., "Housing price prediction via improved ML," *Proc. CS*, 2020.
21. C. Fan et al., "House prices prediction with ML," *Proc. ACM*, 2022.
22. H. Kim et al., "Impact of public rental housing on prices using LSTM," *Sustainability*, 2020.
23. K. Nunna et al., "Time series forecasting of US HPI using ML," *IEEE CSDE*, 2023.
24. S. Shi, "Comparison of real estate price prediction based on LSTM and LGBM," *HSET*, 2023.

25. J. Guo et al., “Can ML improve housing price prediction performance?” *IJSPM*, 2020.
26. M. Yazdani, “ML for real estate price prediction,” arXiv:2110.07151, 2021.
27. C. Lee, “Forecasting spatially correlated targets,” *IJSPM*, 2022.
28. K. Samal et al., “Multi-output spatiotemporal pollution forecasting,” *ASC*, 2022.
29. H. Zhang et al., “House price prediction with textual description data,” *NLE*, 2024.
30. R. Hyndman, *Forecasting: Principles and Practice*, 2021.